✓ **Congratulations! You passed!**
TO PASS 80% or higher

Keep Learning

GRADE
100%

Practical aspects of deep learning

LATEST SUBMISSION GRADE
100%

1. If you have 10,000,000 examples, how would you split the train/dev/test set? — 1 / 1 point

- ○ 60% train . 20% dev . 20% test
- ○ 33% train . 33% dev . 33% test
- ◉ 98% train . 1% dev . 1% test

✓ Correct

2. The dev and test set should: — 1 / 1 point

- ◉ Come from the same distribution
- ○ Come from different distributions
- ○ Be identical to each other (same (x,y) pairs)
- ○ **Have the same number of examples**

✓ Correct

3. If your Neural Network model seems to have high variance, what of the following would be promising things to try? — 1 / 1 point

- ☐ Make the Neural Network deeper
- ☐ Get more test data
- ☐ Increase the number of units in each hidden layer
- ☑ Add regularization

✓ Correct

- ☑ Get more training data

✓ Correct

4. You are working on an automated check-out kiosk for a supermarket, and are building a classifier for apples, bananas and oranges. Suppose your classifier obtains a training set error of 0.5%, and a dev set error of 7%. Which of the following are promising things to try to improve your classifier? (Check all that apply.) — 1 / 1 point

- ☑ Increase the regularization parameter lambda

✓ Correct

- ☐ Decrease the regularization parameter lambda
- ☑ Get more training data

✓ Correct

- ☐ Use a bigger neural network

5. What is weight decay? — 1 / 1 point

- ○ Gradual corruption of the weights in the neural network if it is trained on noisy data.
- ○ The process of gradually decreasing the learning rate during training.
- ○ A technique to avoid vanishing gradient by imposing a ceiling on the values of the weights.
- ◉ A regularization technique (such as L2 regularization) that results in gradient descent shrinking the weights on every iteration.

✓ Correct

6. What happens when you increase the regularization hyperparameter lambda? — 1 / 1 point

- ◉ Weights are pushed toward becoming smaller (closer to 0)
- ○ Weights are pushed toward becoming bigger (further from 0)
- ○ Doubling lambda should roughly result in doubling the weights
- ○ Gradient descent taking bigger steps with each iteration (proportional to lambda)

✓ Correct

7. With the inverted dropout technique, at test time: — 1 / 1 point

- ○ You do not apply dropout (do not randomly eliminate units), but keep the 1/keep_prob factor in the calculations used in training.
- ○ You apply dropout (randomly eliminating units) and do not keep the 1/keep_prob factor in the calculations used in training
- ◉ You do not apply dropout (do not randomly eliminate units) and do not keep the 1/keep_prob factor in the calculations used in training
- ○ You apply dropout (randomly eliminating units) but keep the 1/keep_prob factor in the calculations used in training.

✓ Correct

8. Increasing the parameter keep_prob from (say) 0.5 to 0.6 will likely cause the following: (Check the two that apply) — 1 / 1 point

- ☐ Increasing the regularization effect
- ☑ Reducing the regularization effect

✓ Correct

- ☐ Causing the neural network to end up with a higher training set error
- ☑ Causing the neural network to end up with a lower training set error

✓ Correct

9. Which of these techniques are useful for reducing variance (reducing overfitting)? (Check all that apply.) — 1 / 1 point

- ☐ Exploding gradient
- ☐ Vanishing gradient
- ☑ Data augmentation

✓ Correct

- ☐ Xavier initialization
- ☐ Gradient Checking
- ☑ Dropout

✓ Correct

- ☑ L2 regularization

✓ Correct

10. Why do we normalize the inputs x? — 1 / 1 point

- ○ It makes the parameter initialization faster
- ◉ It makes the cost function faster to optimize
- ○ It makes it easier to visualize the data
- ○ Normalization is another word for regularization--it helps to reduce variance

✓ Correct