

←

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

✓ Congratulations! You passed!

TO PASS 80% or higher

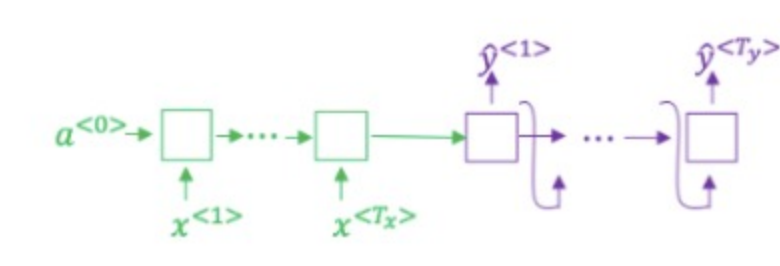
Keep Learning

GRADE
100%

Sequence models & Attention mechanism

LATEST SUBMISSION GRADE
100%

1. Consider using this encoder-decoder model for machine translation.



This model is a "conditional language model" in the sense that the encoder portion (shown in green) is modeling the probability of the input sentence x .

☐ True

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

←

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

✓ Correct

2. In beam search, if you increase the beam width B , which of the following would you expect to be true? Check all that apply.

☒ Beam search will run more slowly.

✓ Correct

☒ Beam search will use up more memory.

✓ Correct

☒ Beam search will generally find better solutions (i.e. do a better job maximizing $P(y \mid x)$)

✓ Correct

☐ Beam search will converge after fewer steps.

3. In machine translation, if we carry out beam search without using sentence normalization, the algorithm will tend to output overly short translations.

☒ True

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

←

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

✓ Correct

4. Suppose you are building a speech recognition system, which uses an RNN model to map from audio clip x to a text transcript y . Your algorithm uses beam search to try to find the value of y that maximizes $P(y \mid x)$.

On a dev set example, given an input audio clip, your algorithm outputs the transcript \hat{y} = "I'm building an A Eye system in Silly con Valley", whereas a human gives a much superior transcript y^* = "I'm building an AI system in Silicon Valley."

According to your model,

$$P(\hat{y} \mid x) = 1.09 \times 10^{-7}$$

$$P(y^* \mid x) = 7.21 \times 10^{-8}$$

Would you expect increasing the beam width B to help correct this example?

☒ No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.

☐ No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.

☐ Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.

☐ Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.

✓ Correct

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

←

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

One way of characterizing the search algorithm is to consider maximizing $P(y \mid x)$ by finding the y that maximizes $P(y \mid x)$. This suggests you should focus your attention on improving the search algorithm.

☒ True.

☐ False.

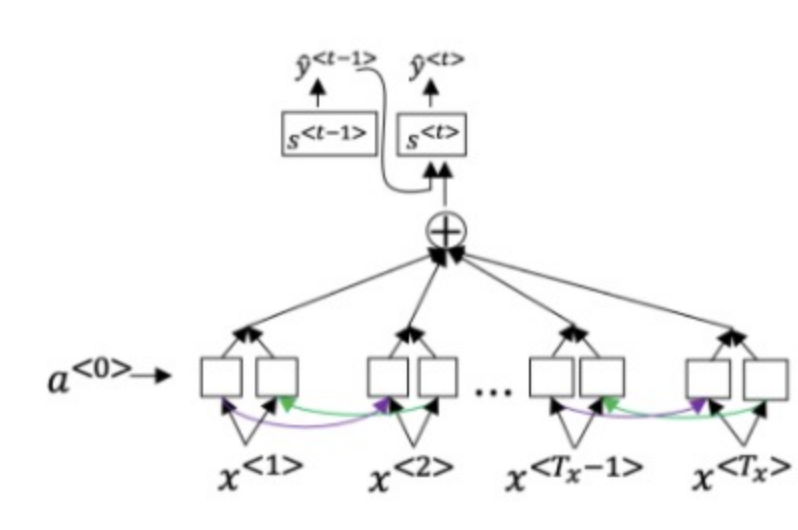
✓ Correct

←

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

6. Consider the attention model for machine translation.



Further, here is the formula for $\alpha^{<t,t'>}$.

$$\alpha^{<t,t'>} = \frac{\exp(e^{<t,t'>})}{\sum_{t'} \exp(e^{<t,t'>})}$$

Which of the following statements about $\alpha^{<t,t'>}$ are true? Check all that apply.

☒ We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t'>}$ that are highly relevant to the value the network should output for $y^{<t'>}$. (Note the indices in the superscripts.)

✓ Correct

☐ We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t'>}$ that are highly relevant to the value the network should output for $y^{<t'>}$. (Note the indices in the superscripts.)

☐ $\sum_{t'} \alpha^{<t,t'>} = 1$ (Note the summation is over t .)

☒ $\sum_{t'} \alpha^{<t,t'>} = 1$ (Note the summation is over t' .)

✓ Correct

7. The network learns where to "pay attention" by learning the values $e^{<t,t'>}$, which are computed using a small neural network:

We can't replace $s^{<t-1>}$ with $s^{<t>}$ as an input to this neural network. This is because $s^{<t>}$ depends on $\alpha^{<t,t'>}$ which in turn depends on $e^{<t,t'>}$; so at the time we need to evaluate this network, we haven't computed $s^{<t>}$ yet.

☒ True

☐ False

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

←

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

8. Compared to the encoder-decoder model shown in Question 1 of this quiz (which does not use an attention mechanism).

☐ The input sequence length T_x is small.

✓ Correct

9. Under the CTC model, identical repeated characters not separated by the "blank" character ($_$) are collapsed. Under the CTC model, what does the following string collapse to?

$$_c_oo_o_kk_b_ooooo_oo_kkk$$

☐ cokbok

☒ cookbook

☐ cook book

☐ cooockbooooookkk

✓ Correct

10. In trigger word detection, $x^{<t>}$ is:

☒ Features of the audio (such as spectrogram features) at time t .

☐ The t -th input word, represented as either a one-hot vector or a word embedding.

☐ Whether the trigger word is being said at time t .

☐ Whether someone has just finished saying the trigger word at time t .

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

←

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST

8. Compared to the encoder-decoder model shown in Question 1 of this quiz (which does not use an attention mechanism).

☐ The input sequence length T_x is small.

✓ Correct

9. Under the CTC model, identical repeated characters not separated by the "blank" character ($_$) are collapsed. Under the CTC model, what does the following string collapse to?

$$_c_oo_o_kk_b_ooooo_oo_kkk$$

☐ cokbok

☒ cookbook

☐ cook book

☐ cooockbooooookkk

✓ Correct

10. In trigger word detection, $x^{<t>}$ is:

☒ Features of the audio (such as spectrogram features) at time t .

☐ The t -th input word, represented as either a one-hot vector or a word embedding.

☐ Whether the trigger word is being said at time t .

☐ Whether someone has just finished saying the trigger word at time t .

Sequence models & Attention mechanism
Graded Quiz • 30 min

Due Feb 17, 3:59 PM CST